

Table 1: Deployment efficiency for serving four requested horizons. Values are macro-averaged over seven datasets on one exclusive RTX 3090 with FP32 and batch size 1.

System	Models	Params (M)	Ckpt. (MiB)	Train (GPU h)	Single (ms)	All- H (ms)	Peak (MiB)	CHPC
ISCF-BSCA	1	2.926	17.68	2.028	10.306	10.318	38.8	Architectural
TimeAlign	4	10.741	95.43	0.290	0.928	3.582	109.1	No
QDF	4	5.337	20.38	0.204	1.141	4.426	31.9	No
DLinear- $H720$ -prefix	1	0.485	1.85	0.021	0.405	0.429	11.1	Service protocol
PatchTST- $H720$ -prefix	1	3.198	12.23	1.110	3.933	3.882	58.0	Service protocol

Models, parameters, checkpoints and training time count the complete deployed service per dataset. Single is the arithmetic mean over requests $H \in \{96, 192, 336, 720\}$. All- H uses one $H = 720$ forward plus prefix views for one-model services and four sequential native forwards for horizon-specific families. Inputs are synthetic standardized tensors; no test loader or labels are accessed.

Train (GPU h) is the sum of native per-epoch training times recorded in the frozen logs and excludes validation, testing and unlogged orchestration overhead.