Table 1: Accuracy, peak inference memory, and checkpoint storage for serving four forecasting horizons. Accuracy is the macro average of the 28 Main I dataset–horizon cells. Memory and storage are macro-averaged over seven datasets. Lower is better.

System	MSE	MAE	Peak memory (MiB)	Checkpoints (MiB)
ISCF-BSCA (ours)	**0.261**	**0.306**	38.817	<u>17.677</u>
TimeAlign	<u>0.274</u>	<u>0.314</u>	109.102	95.433
QDF	0.288	0.331	<u>31.939</u>	20.381
AMD	0.282	0.328	238.035	221.150
SimpleTM	0.293	0.333	**14.841**	**2.835**

ISCF-BSCA stores and serves one unified checkpoint. TimeAlign, QDF, AMD, and SimpleTM sum the four native checkpoints for $H \in \{96, 192, 336, 720\}$. Peak memory is measured on an exclusive RTX 3090 in FP32 with batch size 1 and all service checkpoints resident: one $H = 720$ forward plus prefix views for ISCF-BSCA, and four sequential native forwards for each horizon-specific family. SimpleTM resource cost uses repeat 0 as one deployable, non-ensemble instance per horizon, whereas its frozen Main I accuracy follows the table's multi-run summary. Synthetic standardized inputs are used; no test loader or labels are accessed. Best results are bold and second-best results are underlined.