

Table 1: Accuracy, peak inference memory, and checkpoint storage for serving four forecasting horizons. Accuracy is the macro average of the 28 Main I dataset–horizon cells. Memory and storage are macro-averaged over seven datasets. Lower is better.

System	MSE	MAE	Peak memory (MiB)	Checkpoints (MiB)
ISCF-BSCA (ours)	0.261	0.306	38.817	17.677
TimeAlign	<u>0.274</u>	<u>0.314</u>	109.102	95.433
QDF	0.288	0.331	31.939	20.381
AMD	0.282	0.328	238.035	221.150
SimpleTM	0.293	0.333	<u>14.841</u>	2.835
DLinear	0.300	0.339	13.111	<u>3.465</u>
iTransformer	0.292	0.336	47.341	35.425
PatchTST	0.280	0.326	71.604	25.295
TimeMixer	0.283	0.332	53.422	37.182

ISCF-BSCA stores and serves one unified checkpoint. Each baseline sums four horizon-specific models for $H \in \{96, 192, 336, 720\}$. Peak memory is measured on an exclusive RTX 3090 in FP32 with batch size 1 and all service models resident: one $H = 720$ forward plus prefix views for ISCF-BSCA, and four sequential native forwards for each baseline. TimeAlign, QDF, AMD, and SimpleTM use actual trained checkpoint files (SimpleTM repeat 0); DLinear, iTransformer, PatchTST, and TimeMixer use official-configuration architecture-equivalent serialized state dicts because a complete local 7×4 trained-checkpoint inventory is unavailable. Synthetic standardized inputs are used; no test loader or labels are accessed. Best results are bold and second-best results are underlined.