

Table 1: Accuracy and training-resource cost for serving four forecasting horizons. Accuracy is the macro average of the 28 Main I dataset-horizon cells. Parameter count and one-epoch cycle time are macro-averaged over seven datasets; horizon-specific systems sum their four native models. Lower is better.

System	Models	MSE	MAE	Params (M)	1-Epoch (s)
ISCF-BSCA (ours)	1	0.261	0.306	2.926	133.1
TimeAlign	4	0.274	0.314	10.741	104.6
QDF	4	0.288	0.331	5.337	55.4

ISCF-BSCA counts one unified model and one training epoch. TimeAlign and QDF count the sum of four independently trained models for $H \in \{96, 192, 336, 720\}$. For each checkpoint, 1-Epoch is the median logged duration of a native training epoch plus its scheduled validation pass; no test evaluation is included.